

**Visa**

**Big Data 101**

**Student Lab 4**



Web Age Solutions Inc.  
USA: 1-877-517-6540  
Canada: 1-866-206-4644  
Web: <http://www.webagesolutions.com>

The following terms are trademarks of other companies:

Java and all Java-based trademarks and logos are trademarks or registered trademarks of Sun Microsystems, Inc. in the United States, other countries, or both.

Microsoft, Windows, Windows NT, and the Windows logo are trademarks of Microsoft Corporation in the United States, other countries, or both.

UNIX is a registered trademark of The Open Group in the United States and other countries.

Linux is a trademark of Linus Torvalds in the United States, other countries, or both.

IBM, WebSphere, DB2 and Tivoli are trademarks of the International Business Machines

Corporation in the United States, other countries, or both.

Other company, product, and service names may be trademarks or service marks of others.

For customizations of this book or other sales inquiries, please contact us  
at:

USA: 1-877-517-6540, email: [getinfousa@webagesolutions.com](mailto:getinfousa@webagesolutions.com)

Canada: 1-866-206-4644 toll free, email: [getinfo@webagesolutions.com](mailto:getinfo@webagesolutions.com)

Copyright © 2022 Web Age Solutions Inc.

This publication is protected by the copyright laws of Canada, United States and any other country where this book is sold. Unauthorized use of this material, including but not limited to, reproduction of the whole or part of the content, re-sale or transmission through fax, photocopy or e-mail is prohibited. To obtain authorization for any such activities, please write to:

Web Age Solutions Inc.  
439 University Ave  
Suite 820  
Toronto  
Ontario, M5G 1Y8

## Contents

<b>Lab 4 –HiveQL Joins and Mini-Project - 1 ½ hours</b> .....	4
<b>Step 1.</b> Connect to our Lab Environment .....	4
<b>Step 2.</b> Joins.....	5
<b>Step 3.</b> Inner Joins .....	7
<b>Step 4.</b> Left Outer Join.....	10
<b>Step 5.</b> Additional Join Optimization .....	11
<b>Mini-Project</b> .....	14
<b>Review</b> .....	19

## Lab 4 –HiveQL Joins and Mini-Project - 1 ½ hours

### Overview

### HiveQL Specialized Queries and Mini-Project

In this lab we start off with some new ways of working with variables in the Hive Environment. We will then work with the Hive CLI and apply what we have already learned throughout the class to solve various issues with the data and to build upon our reporting skillset.

#### Step 1. Connect to our Lab Environment

Please connect to the assigned Nuvepro Virtual Machine where the labs will be run from.

On the Linux desktop you will see the HDP VM has already been started and is showing a couple of ways to login to the system.

Open a web browser on the Nuvepro Virtual Machine. Enter this address into the URL to open the SSH connection to HortonWorks HDP Sandbox.

<http://localhost:4200>



Login to the console and provide the password when prompted.

Login username: **wasadmin**

Password: **wasadmin**

We will now list the directory information using the Linux `ll` command (lowercase letter L) for long listing. You should see LabFiles folder.

**ll**

Use the following command to change to that directory.

**cd LabFiles**

Start the Hive CLI

**hive**

Set the database to use wasadmin\_db

**use wasadmin\_db;**

Make sure the tables are in the database.

**show tables;**

## Step 2. Joins

The Hive QL syntax for joining is like SQL with a few added conditional operations like MAPJOIN (bucketmapjoin, sortedmerge) and enforcement of parenthesis in the on condition syntax.

If all but one of the tables being joined are small, the join can be performed as a map only job saving resources and time.

<https://cwiki.apache.org/confluence/display/hive/languagemanual+joins>

Hive supports the following syntax for joining tables:

join\_table:

table\_reference [INNER] JOIN table\_factor [join\_condition]

```

| table_reference {LEFT|RIGHT|FULL} [OUTER] JOIN table_reference join_condition
| table_reference LEFT SEMI JOIN table_reference join_condition
| table_reference CROSS JOIN table_reference [join_condition] (as of Hive 0.10)

```

table\_reference:

```

    table_factor
| join_table

```

table\_factor:

```

    tbl_name [alias]
| table_subquery alias
| ( table_references )

```

join\_condition:

```

    ON expression

```

To combine and retrieve the records from multiple tables we use a Hive Join clause taking into consideration how the tables are related. That means knowing the many data tables' columns and the commonality between the tables to join is a must. First let's look at the tables we want to join and find the common columns.

Start by using the describe command. Here we will also look at additional information on both transactions and vendors tables.

## describe extended transactions;

```

+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| id | int | |
| account_id | int | |
| vendor_id | int | |
| times | string | |
| city | string | |
| state | string | |
| amount | double | |
| | NULL | NULL |
| Detailed Table Information | Table(tableName:transactions, dbName:wasadmin_db, owner:hive, createTime:1642641890, lastAccessTime:0, retention:0, sd:StorageDescriptor(cols:[FieldSchema(name:id, type:int, comment:null), FieldSchema(name:account_id, type:int, comment:null), FieldSchema(name:vendor_id, type:int, comment:null), FieldSchema(name:times, type:string, comment:null), FieldSchema(name:city, type:string, comment:null), FieldSchema(name:state, type:string, comment:null), FieldSchema(name:amount, type:double, comment:null)], location:hdfs://sandbox-hdp.hortonworks.com:8020/data/transactions/in, inputFormat:org.apache.hadoop.mapred.TextInputFormat, outputFormat:org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed:false, numBuckets:-1, serDeInfo:SerDeInfo(name:null, serializationLib:org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters:{serialization.format=, field.delim=}), bucketCols:[], sortCols:[], parameters:{}, skewedInfo:SkewedInfo(skewedColNames:[], skewedColValues:[], skewedColValueLocationMaps:{}), storedAsSubDirectories:false), partitionKeys:[], parameters:{transient_lastDdlTime=1642641890, bucketing_version=2, totalSize=48168931, EXTERNAL=TRUE, numFiles=11}, viewOriginalText:null, viewExpandedText:null, tableType:EXTERNAL_TABLE, rewriteEnabled:false, catName:hive, ownerType:USER, writeId:0) |
+-----+-----+-----+
9 rows selected (0.347 seconds)

```

**describe extended vendors;**

```
+-----+-----+-----+
| col_name | data_type | comment |
+-----+-----+-----+
| id        | int       |         |
| name      | string    |         |
| city      | string    |         |
| state     | string    |         |
| category  | string    |         |
| swipe_rate | double    |         |
| NULL      | NULL      | NULL    |
+-----+-----+-----+
| Detailed Table Information | Table(tableName=vendors, dbName=wasadmin_db, owner=hive, createTime=1642641500, lastAccessTime=0, retention=0, sd=StorageDescriptor(cols=[FieldSchema(name=id, type=int, comment=null), FieldSchema(name=name, type:string, comment=null), FieldSchema(name=city, type:string, comment=null), FieldSchema(name=state, type:string, comment=null), FieldSchema(name=category, type:string, comment=null), FieldSchema(name=swipe_rate, type=double, comment=null)], location:hdfs://sandbox-hdp.hortonworks.com:8020/data/vendors/in, inputFormat=org.apache.hadoop.mapred.TextInputFormat, outputFormat=org.apache.hadoop.hive.ql.io.HiveIgnoreKeyTextOutputFormat, compressed=false, numBuckets=-1, serdeInfo=SerDeInfo(name=null, serializationLib=org.apache.hadoop.hive.serde2.lazy.LazySimpleSerDe, parameters={serialization.format=, field.delim=,}), bucketCols=[], sortCols=[], parameters={}), skewedInfo=SkewedInfo(skewedColNames=[], skewedColValues=[], skewedColValueLocations=[]), storedAsSubDirectories=false), partitionKeys=[], parameters={transient_lastDdlTime=1642641500, bucketing_version=2, totalSize=371, EXTERNAL=TRUE, numFiles=1}, viewOriginalText=null, viewExpandedText=null, tableType=EXTERNAL_TABLE, rewriteEnabled=false, catName=hive, ownerType=USER, writeId=0) |
+-----+-----+-----+
8 rows selected (0.125 seconds)
```

In the description of the columns, you can see there is an id column in vendors table and a vendor\_id column in the transactions table. While not using the same name for the columns these are the common columns between these tables.

**Q: How could we confirm that?**

**A: Check the data in those fields by writing simple queries.**

**select \* from transactions limit 10;**

**select \* from vendors limit 10;**

### Step 3. Inner Joins

Time to try our command to retrieve the columns and common data between the two tables. In this case we must set the on condition. Here we join 'transactions' and 'vendors' tables.

**select transactions.\*, vendors.\* from transactions join vendors on (transactions.vendor\_id = vendors.id) limit 10;**

Generally, it is better to return a projection of columns when querying data. That means selecting only specific columns to be returned using comma separation between column names.

Find the sum of all amounts by account\_id:

```
select transactions.account_id, SUM(amount) as total from  
transactions join vendors on (transactions.vendor_id = vendors.id)  
group by account_id limit 10;
```

Did you notice the extra reducer required for the group by aggregation?

VERTICES	MODE	STATUS	TOTAL	COMPLETED	RUNNING	PENDING	FAILED	KILLED
Map 3 .....	container	SUCCEEDED	1	1	0	0	0	0
Map 1 .....	container	SUCCEEDED	4	4	0	0	0	0
Reducer 2 .....	container	SUCCEEDED	2	2	0	0	0	0
VERTICES: 03/03 [=====>>] 100% ELAPSED TIME: 34.89 s								

Here the columns are easier to read since we have chosen a smaller number to be returned (projected) as the resultset.

transactions.account_id	total
1	13285.3
2	13739.179999999998
3	11646.52
4	10724.279999999999
5	16656.16
8	12592.369999999999
10	14929.560000000001
6	13142.0
7	13603.510000000002
9	15849.209999999997

10 rows selected (20.074 seconds)

**Q: How would you find the sum of all amounts by vendors name??**

**A: Add the vendors.name column to the projection, and don't forget to group by it too.**



transactions.account_id	vendors.name	total
1	American Airlines	489.46
1	Apple	1531.1800000000003
1	Best Buy	1022.03
1	Kohls	1321.89
1	REI	1520.6000000000001
2	American Airlines	521.6299999999999
2	Apple	993.39
2	Barnes and Noble	195.06
2	REI	14.86
4	Starbucks	42.0

10 rows selected (15.957 seconds)

**select transactions.account\_id, vendors.name, SUM(amount) as total from transactions join vendors on (transactions.vendor\_id = vendors.id) group by account\_id, vendors.name limit 10;**

**Challenge 1:** (Solutions are not shown, confer with others if stuck on a question)

**Find the discount offered for the top 5 transactions (highest spent).  
Hint: Join offers and transactions tables.**

**Use the offers table discount field. Calculate the average amount spent on each category. Show the top 5 results with the highest discount.**

**What question does the following query answer?**

**select vendors.category, SUM(transactions.amount) as total from transactions join vendors on (transactions.vendor\_id = vendors.id) group by vendors.category limit 10;**

HiveQL can join multiple tables as required if the configuration within the system allocates enough memory. Like SQL engines, Hive will try to

load the smallest table in memory then use caching and buffering as necessary to work with the additional tables.

**select transactions.\*, vendors.\*, offers.discount as discount from transactions join vendors on (transactions.vendor\_id = vendors.id) join offers on (offers.vendor\_id = vendors.id) limit 10**

transactions.id	transactions.account_id	transactions.vendor_id	transactions.times	transactions.city	transactions.state	transactions.amount	vendors.id	vendors.name	vendors.city		
vendors.state	vendors.category	vendors.suicide_rate	discount								
600001	9476	3.2	8	0.1	2015-01-07 00:00:00	Springfield	MA	94.22	8	Hilton	Washington
600002	9302	3.5	7	0.11	2015-01-07 00:00:00	Bronx	NY	283.87	7	Barnes and Noble	New York
NY	Books	2.5	4	0.1	2015-01-07 00:00:02	Savannah	GA	195.83	4	Apple	Cupertino
600004	7335	2.5	4	0.1	2015-01-07 00:00:04	New Orleans	LA	56.05	4	Apple	Cupertino
CA	Technology	3.0	5	0.1	2015-01-07 00:00:05	Springfield	MA	24.22	5	Best Buy	Sacramento
600006	2594	2.4	6	0.12	2015-01-07 00:00:06	Baltimore	MD	253.21	6	Kohls	Minneapolis
CA	Technology	2.0	1	0.05	2015-01-07 00:00:08	San Francisco	CA	39.67	1	Walmart	Bentonville
600007	947	3.5	7	0.11	2015-01-07 00:00:09	Indianapolis	IN	139.96	7	Barnes and Noble	New York
CA	Technology	3.2	8	0.1	2015-01-07 00:00:10	Springfield	MA	23.98	8	Hilton	Washington
600009	9521	3.2	8	0.1	2015-01-07 00:00:11	Peoria	IL	19.98	8	Hilton	Washington
MN	Fashion										
600011	7081										
AR	General										
600012	8793										
NY	Books										
600013	328										
DC	Travel										
600014	1340										
DC	Travel										

10 rows selected (1390.594 seconds)

10 rows selected (1390.594 seconds)

Without using partitioning or bucketing, Hive reads all the data in the directory and applies the query filters on it. This is *slow and expensive* for memory since all data must be read in each of the three tables.

```
INFO : PHYSICAL_MEMORY_BYTES: 2943352832
INFO : VIRTUAL_MEMORY_BYTES: 15919079424
INFO : COMMITTED_HEAP_BYTES: 2943352832
```

## Step 4. Left Outer Join

The HiveQL LEFT OUTER JOIN returns all the rows from the left table, even if there are no matches in the right table. A LEFT JOIN returns all the values from the left table, plus the matched values from the right table, or NULL in case of no matching JOIN predicate.

<https://cwiki.apache.org/confluence/display/Hive/OuterJoinBehavior>

The following query demonstrates Left Outer Join between transactions and offers tables, notice the result shows null values.

**select transactions.\*, offers.\* from transactions [left outer join](#) offers on (transactions.vendor\_id = offers.vendor\_id) limit 10;**

**Challenge 2:** (Solutions are not shown, confer with others if stuck on a question)

**Picking two tables at a time identify any missing information when joining these tables; Accounts, Offers, Transactions and Vendors.**

## **Right Outer Join**

A right outer join behaves the opposite to a left outer join. In this join, all the rows from the right table and the matching rows from the left table are displayed. All the unmatched rows from the table on the left will be dropped and null will be displayed.

## **Full Outer Join**

The major purpose of the HiveQL Full Outer Join is it combines the records of both the left and the right outer tables which fulfills the Hive JOIN condition. Moreover, this joined table contains either all the records from both the tables, or fills in NULL values for missing matches on either side.

## **[Step 5. Additional Join Optimization](#)**

Hive automatically recognizes various use cases and optimizes for them. Hive 0.11 improves the optimizer for these cases:

- Joins where one side fits in memory. In the new optimization:
  - that side is loaded into memory as a hash table
  - only the larger table needs to be scanned
  - fact tables have a smaller footprint in memory
- Star-schema joins
- Hints are no longer needed for many cases.
- Map joins are automatically picked up by the optimizer.

Review the documentation here:

<https://cwiki.apache.org/confluence/display/Hive/LanguageManual+JoinOptimization>

## Star Join Optimization

A simple schema for decision support systems or data warehouses is the star schema, where events are collected in large *fact tables*, while smaller supporting tables (*dimensions*) are used to describe the data.

Hive supports MAPJOINS, which are well suited for this scenario – at least for dimensions small enough to fit in memory (less than 25MB). MAPJOINS are processed by loading the smaller table into an in-memory hash map and matching keys with the larger table as they are streamed through.

```
select /*+ MAPJOIN(b) */ a.key, a.value from a join b on a.key = b.key
```

```
select accounts.* from accounts a, offers b where a.account_id = b.account_id;
```

## Challenge 3:

**Calculate total invoice for each account. Hint : group by account\_id**

**Find top 10 spending customers by name. What is the query?**

**Calculate accounts totals by month**

Exit from the hive shell once completed.

**!q**

## Mini-Project

In this Mini-Project you'll learn to work with n-grams on a text file to identify aspects within the file itself. N-grams of texts are extensively used in text mining and natural language processing tasks. They are basically a set of co-occurring words within a given window and when computing the n-grams you typically move one word forward (although you can move X words forward in more advanced scenarios).

<https://en.wikipedia.org/wiki/N-gram>

Hive uses command formula to create the ngrams in two ways; ngrams and context\_ngrams. We will look at both in the Mini-Project.

```
SELECT ngrams(sentences(lower(tweet)), 2, 100 [, 1000]) FROM twitter;
```

```
SELECT context_ngrams(sentences(lower(tweet)), array("i", "love", null), 100, [, 1000]) FROM twitter;
```

For more information see:

[https://cwiki.apache.org/confluence/display/hive/statisticsanddatamining#StatisticsAndDataMining-ngrams\(\)andcontext\\_ngrams\(\):N-gramfrequencyestimation](https://cwiki.apache.org/confluence/display/hive/statisticsanddatamining#StatisticsAndDataMining-ngrams()andcontext_ngrams():N-gramfrequencyestimation)

We will start by finding the file within the HDFS file system. The file is called *moby-dick.txt*, the classic novel.

1. Use Hadoop CLI commands to find the file and the folder path or browse the file system using the Browser Utility. Make a note of the path.
2. Start Hive CLI and switch to the wasadmin\_db database.
3. Check to see if the table the 'moby' table already exists.
4. If the table exists use it, otherwise create the external table with the following properties:

- a. Column name line, data type is string
  - b. Rows are delimited
  - c. All fields have a newline character at the end
  - d. It is stored as a text file
  - e. Found in the path you identified within HDFS
5. Ensure the data from the moby-dick.txt file was loaded properly into the hive 'moby' table you created. Examine the data in the table.
  6. Find the 20 most occurring ngrams (one) word in the text.  
<https://en.wikipedia.org/wiki/N-gram>. In Hive we can explode, format the output of, the 'ngrams' by reading in each line of text from the file and applying the value of how many, in this case 1)

You should be able to produce this output:

col
{"ngram":["the"],"estfrequency":14617.0}
{"ngram":["of"],"estfrequency":6709.0}
{"ngram":["and"],"estfrequency":6488.0}
{"ngram":["a"],"estfrequency":4761.0}
{"ngram":["to"],"estfrequency":4678.0}
{"ngram":["in"],"estfrequency":4223.0}
{"ngram":["that"],"estfrequency":3005.0}
{"ngram":["his"],"estfrequency":2530.0}
{"ngram":["it"],"estfrequency":2434.0}
{"ngram":["i"],"estfrequency":1989.0}
{"ngram":["but"],"estfrequency":1823.0}
{"ngram":["he"],"estfrequency":1780.0}
{"ngram":["with"],"estfrequency":1770.0}
{"ngram":["as"],"estfrequency":1751.0}
{"ngram":["is"],"estfrequency":1749.0}
{"ngram":["for"],"estfrequency":1646.0}
{"ngram":["was"],"estfrequency":1646.0}
{"ngram":["all"],"estfrequency":1527.0}
{"ngram":["this"],"estfrequency":1443.0}
{"ngram":["at"],"estfrequency":1334.0}

7. Find the 20 most occurring bigrams (two) words in the text.  
<https://en.wikipedia.org/wiki/Bigram>. In Hive we can explode, format the output of, the 'ngrams' by reading in each line of text from the file and applying the value of how many, in this case 2)  
 You should be able to produce this output:

col
{ "ngram": ["of", "the"], "estfrequency": 1796.0 }
{ "ngram": ["in", "the"], "estfrequency": 1144.0 }
{ "ngram": ["to", "the"], "estfrequency": 708.0 }
{ "ngram": ["from", "the"], "estfrequency": 408.0 }
{ "ngram": ["of", "his"], "estfrequency": 359.0 }
{ "ngram": ["and", "the"], "estfrequency": 351.0 }
{ "ngram": ["on", "the"], "estfrequency": 342.0 }
{ "ngram": ["to", "be"], "estfrequency": 322.0 }
{ "ngram": ["of", "a"], "estfrequency": 319.0 }
{ "ngram": ["at", "the"], "estfrequency": 317.0 }
{ "ngram": ["by", "the"], "estfrequency": 307.0 }
{ "ngram": ["with", "the"], "estfrequency": 302.0 }
{ "ngram": ["for", "the"], "estfrequency": 297.0 }
{ "ngram": ["the", "whale"], "estfrequency": 292.0 }
{ "ngram": ["it", "was"], "estfrequency": 278.0 }
{ "ngram": ["it", "is"], "estfrequency": 269.0 }
{ "ngram": ["in", "his"], "estfrequency": 255.0 }
{ "ngram": ["in", "a"], "estfrequency": 254.0 }
{ "ngram": ["with", "a"], "estfrequency": 237.0 }
{ "ngram": ["into", "the"], "estfrequency": 236.0 }

8. Find the 20 most occurring trigrams (three) words in the text.  
<https://en.wikipedia.org/wiki/Trigram>. In Hive we can explode, format the output of, the ‘ngrams’ by reading in each line of text from the file and applying the value of how many, in this case 2)

You should be able to produce this output:

col
{ "ngram": ["of", "the", "whale"], "estfrequency": 67.0 }
{ "ngram": ["the", "sperm", "whale"], "estfrequency": 67.0 }
{ "ngram": ["the", "white", "whale"], "estfrequency": 67.0 }
{ "ngram": ["of", "the", "sea"], "estfrequency": 50.0 }
{ "ngram": ["part", "of", "the"], "estfrequency": 50.0 }
{ "ngram": ["one", "of", "the"], "estfrequency": 46.0 }
{ "ngram": ["out", "of", "the"], "estfrequency": 44.0 }
{ "ngram": ["a", "sort", "of"], "estfrequency": 35.0 }
{ "ngram": ["of", "the", "sperm"], "estfrequency": 31.0 }
{ "ngram": ["it", "was", "a"], "estfrequency": 30.0 }
{ "ngram": ["in", "the", "sea"], "estfrequency": 28.0 }
{ "ngram": ["it", "is", "a"], "estfrequency": 25.0 }
{ "ngram": ["of", "the", "boat"], "estfrequency": 25.0 }
{ "ngram": ["to", "the", "deck"], "estfrequency": 25.0 }
{ "ngram": ["in", "order", "to"], "estfrequency": 22.0 }
{ "ngram": ["of", "the", "pequod"], "estfrequency": 22.0 }
{ "ngram": ["into", "the", "sea"], "estfrequency": 21.0 }
{ "ngram": ["the", "sea", "and"], "estfrequency": 21.0 }
{ "ngram": ["it", "was", "not"], "estfrequency": 20.0 }
{ "ngram": ["there", "was", "a"], "estfrequency": 20.0 }

9. In the first two attempts we have identified small words often known as “stop words”, and ones we often want to ignore. Using the context\_ngram we can search for phrases like “sperm whale” or “white whale” or just “white” if you want.



White whale:

col
{"ngram":["had"],"estfrequency":7.0}
{"ngram":["a"],"estfrequency":3.0}
{"ngram":["the"],"estfrequency":3.0}
{"ngram":["as"],"estfrequency":2.0}
{"ngram":["did"],"estfrequency":2.0}
{"ngram":["he"],"estfrequency":2.0}
{"ngram":["is"],"estfrequency":2.0}
{"ngram":["must"],"estfrequency":2.0}
{"ngram":["on"],"estfrequency":2.0}
{"ngram":["shirr"],"estfrequency":2.0}
{"ngram":["that"],"estfrequency":2.0}
{"ngram":["was"],"estfrequency":2.0}
{"ngram":["and"],"estfrequency":1.0}
{"ngram":["are"],"estfrequency":1.0}
{"ngram":["be"],"estfrequency":1.0}
{"ngram":["by"],"estfrequency":1.0}
{"ngram":["dashed"],"estfrequency":1.0}
{"ngram":["from"],"estfrequency":1.0}
{"ngram":["so"],"estfrequency":1.0}
{"ngram":["spouts"],"estfrequency":1.0}

Sperm whale:

col
{"ngram":["is"],"estfrequency":10.0}
{"ngram":["and"],"estfrequency":6.0}
{"ngram":["was"],"estfrequency":5.0}
{"ngram":["fishery"],"estfrequency":4.0}
{"ngram":["has"],"estfrequency":4.0}
{"ngram":["in"],"estfrequency":3.0}
{"ngram":["of"],"estfrequency":3.0}
{"ngram":["only"],"estfrequency":3.0}
{"ngram":["will"],"estfrequency":3.0}
{"ngram":["before"],"estfrequency":2.0}
{"ngram":["but"],"estfrequency":2.0}
{"ngram":["for"],"estfrequency":2.0}
{"ngram":["now"],"estfrequency":2.0}
{"ngram":["presents"],"estfrequency":2.0}
{"ngram":["to"],"estfrequency":2.0}
{"ngram":["which"],"estfrequency":2.0}
{"ngram":["after"],"estfrequency":1.0}
{"ngram":["compared"],"estfrequency":1.0}
{"ngram":["inserted"],"estfrequency":1.0}
{"ngram":["on"],"estfrequency":1.0}

White:

col
{"ngram":["whale"],"estfrequency":86.0}
{"ngram":["whale's"],"estfrequency":10.0}
{"ngram":["water"],"estfrequency":7.0}
{"ngram":["man"],"estfrequency":5.0}
{"ngram":["hump"],"estfrequency":4.0}
{"ngram":["bears"],"estfrequency":3.0}
{"ngram":["mass"],"estfrequency":3.0}
{"ngram":["all"],"estfrequency":2.0}
{"ngram":["and"],"estfrequency":2.0}
{"ngram":["as"],"estfrequency":2.0}
{"ngram":["bubbles"],"estfrequency":2.0}
{"ngram":["cedar"],"estfrequency":2.0}
{"ngram":["flag"],"estfrequency":2.0}
{"ngram":["for"],"estfrequency":2.0}
{"ngram":["head"],"estfrequency":2.0}
{"ngram":["one"],"estfrequency":2.0}
{"ngram":["shark"],"estfrequency":2.0}
{"ngram":["squalls"],"estfrequency":2.0}
{"ngram":["steed"],"estfrequency":2.0}
{"ngram":["vapours"],"estfrequency":2.0}

### -----Bonus-----

There are two additional text files within the HDFS file system. Using similar techniques that you used to work with Moby Dick above apply your understanding to complete the following with these files.

- **sotu-2014-obama.txt**  
State of the Union address 2014

#### **For State of the union text**

Find the word the follows the word 'american'

- **FOMC20080916meeting.txt**  
Feds emergency meeting after 2008 financial crisis (unclassified after 5 years)

#### **Feds Meeting minute**

Find the most used adjective used to describe the word market.

## Review

In this last lab and Mini-Project we have learned about working with tables especially joining them to project answer to specific questions. Remember knowing the data that you are working with is the first step in being successfully in Hive and Hadop.